

Deep Reinforcement Learning with Double Q-Learning

Qing Lyu

11/07/2018

Reinforcement learning by Google DeepMind

P

Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence (AAAI-16)

Vol

Deep Reinforcement Learning with Double Q-Learning

oglou

Hado van Hasselt, Arthur Guez, and David Silver
Google DeepMind

{v1ad Martin Riedmiller¹, Andreas K. Fidjeland¹, Georg Ostrovski¹, Stig Petersen¹, Charles Beattie¹, Amir Sadik¹, Ioannis Antonoglou¹, Helen King¹, Dhharshan Kumaran¹, Daan Wierstra¹, Shane Legg¹ & Demis Hassabis¹

Abstract

We present the first deep learning model to successfully learn control policies directly from high-dimensional sensory input using reinforcement learning. The model is a convolutional neural network, trained with a variant of Q-learning, whose input is raw pixels and whose output is a value function estimating future rewards. We apply our method to seven Atari 2600 games from the Arcade Learning Environment, with no adjustment of the architecture or learning algorithm. We find that it outperforms all previous approaches on six of the games and surpasses a human expert on three of them.

Background

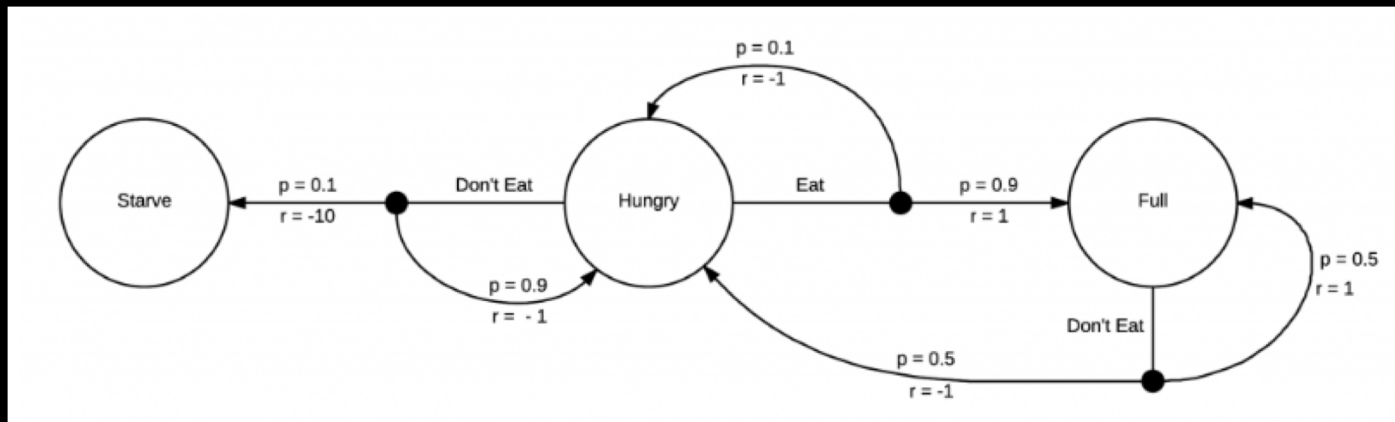
- Cumulative Reward

$$R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

- Policy

- It is a function that takes in a state and an action and returns the probability of taking that action in that state.

$$\sum_a \pi(s, a) = 1$$



Background

- Value Functions
 - **state value function**: the value of a state when following a policy

$$V^{\pi}(s) = \mathbb{E}_{\pi}[R_t | s_t = s]$$

- **action value function**: the value of taking an action in some state when following a certain policy

$$Q^{\pi}(s, a) = \mathbb{E}_{\pi}[R_t | s_t = s, a_t = a]$$

Background

- **Q-learning algorithm**

- Before learning begins, Q is initialized to a possibly arbitrary fixed value (chosen by the programmer). Then, at each time t the agent selects an action a_t , observes a reward r_t , enters a new state s_{t+1} (that may depend on both the previous state s_t and the selected action), and Q is updated. The core of the algorithm is a simple value iteration update, using the weighted average of the old value and the new information:

$$Q(s_t, a_t) \leftarrow (1 - \alpha) \cdot \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \overbrace{\left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} \right)}^{\text{learned value}}$$

Example: flappy bird

Two actions:

Fly upward or do nothing

Reward:

Keep alive: +1

Dead: -1000

Pass through a tube: +50

| state | Fly | Do nothing |
|--------------------|-----|------------|
| (dx_1, dy_1) | 0 | 0 |
| (dx_1, dy_2) | 0 | 0 |
| ... | ... | ... |
| (dx_m, dy_{n-1}) | 0 | 0 |
| (dx_m, dy_n) | 0 | 0 |



Background

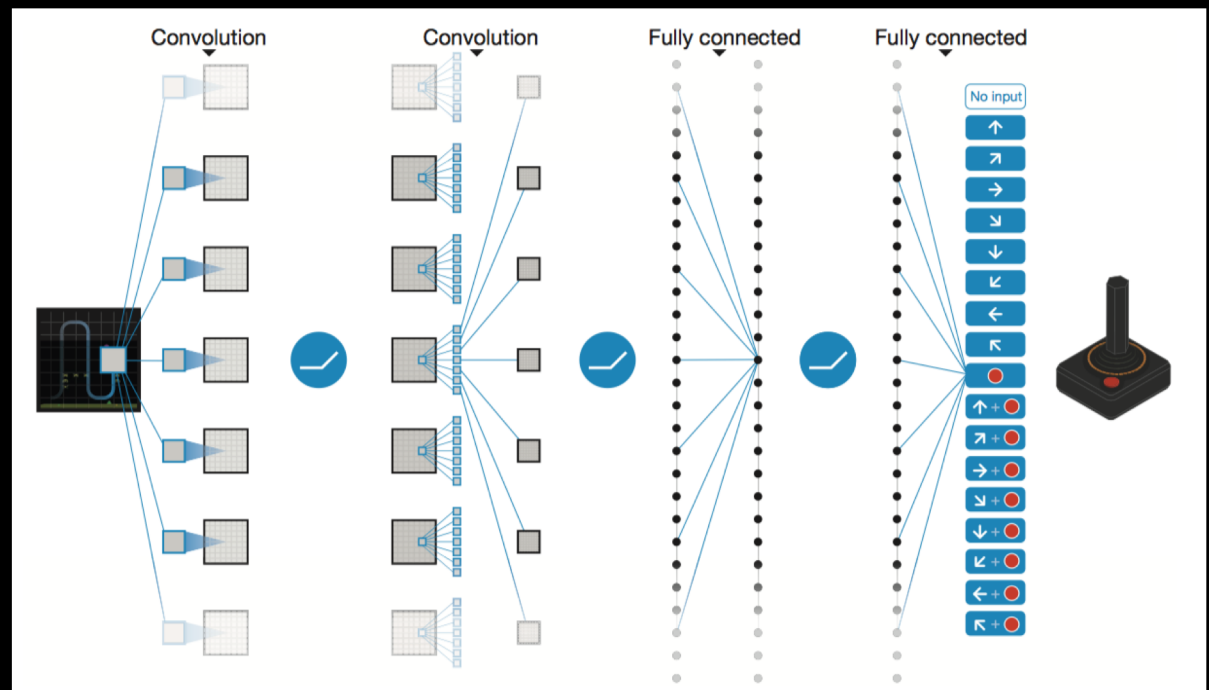
- Deep Q-learning Network
 - target network
 - Experience replay

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

3 convolution layers and 2
fully-connected layers

Input: $84 \times 84 \times 4$ image

Output: a fully-connected
linear layer with a single
output for each valid action



Problems

- Q-learning algorithm tends to overestimate action values under certain conditions
- These overestimations may harm performance

Reason

- The max operator uses the same values both to select and to evaluate an action

$$Y_t^Q \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t)$$

Double Q-learning algorithm


$$Y_t^Q \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t)$$



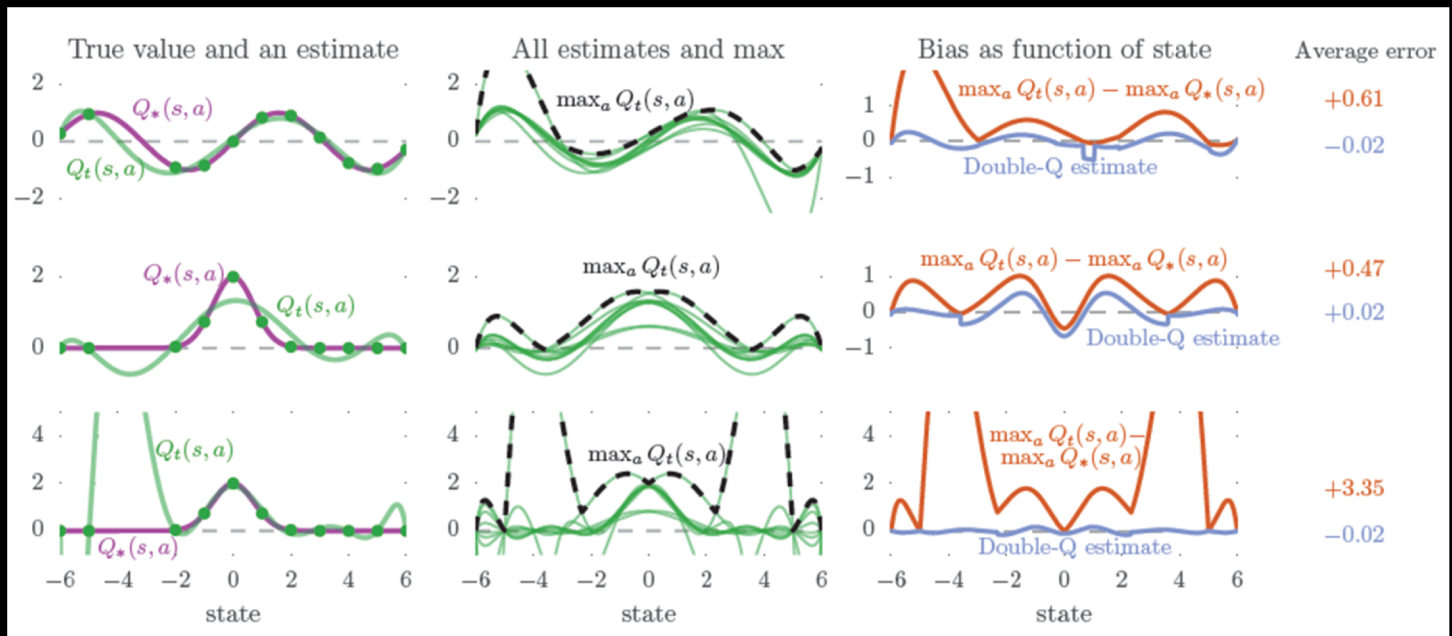
$$Y_t^{\text{DoubleQ}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t); \theta'_t)$$

decomposing the max operation in the target into action selection and action evaluation

Double Q-learning algorithm can decrease overestimation errors



$$Q(s, a) = \sin(s)$$



$$Q(s, a) = 2e^{-s^2}$$


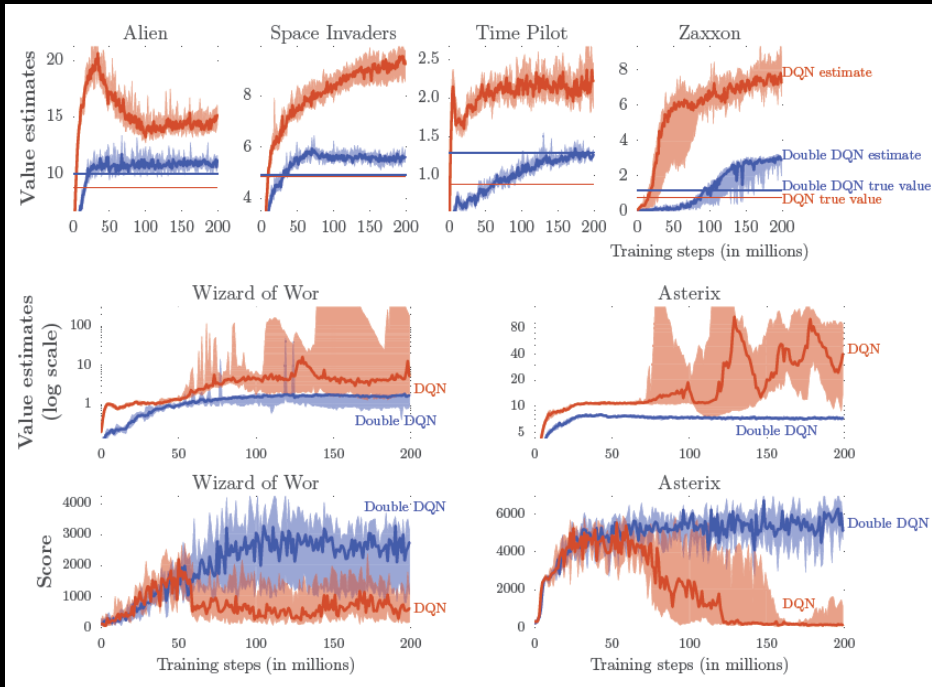
Double Deep Q-learning Network

$$C(\Theta) = \sum_i \left[R^{(i)} + \gamma \max_{a'} f_{Q^*}(s^{(i+1)}, a'; \Theta^-) - f_{Q^*}(s^{(i)}, a^{(i)}; \Theta) \right]^2$$



$$C(\Theta) = \sum_i \left[R^{(i)} + \gamma f_{Q^*}(s^{(i+1)}, \arg \max_{a'} f_{Q^*}(s^{(i+1)}, a'; \Theta); \Theta^-) - f_{Q^*}(s^{(i)}, a^{(i)}; \Theta) \right]^2$$

Result



Contributions

- Shows why Q-learning can be overoptimistic in large-scale problems
- Shows that overestimations are more common and severe in practice than previously acknowledged
- Double Q-learning can be used at scale to successfully reduce this overoptimism, resulting in more stable and reliable learning
- Proposes a Double DQN
- Shows that Double DQN finds better policies, obtaining new state-of-the-art results on the Atari 2600 domain

Related work by Google DeepMind

- Prioritized DDQN <https://arxiv.org/abs/1511.05952>
- Dueling DDQN <https://arxiv.org/abs/1511.06581>
- A3C <https://arxiv.org/abs/1602.01783>
- Distributional DQN <https://arxiv.org/abs/1707.06887>
- Noisy DQN <https://arxiv.org/abs/1706.10295>
- Rainbow <https://arxiv.org/abs/1710.02298>